

Cards to rows: the Oscars, one step at a time

Week 6 · Documents module · Textbook Chapter 6, Strategy 6

Brian C. Keegan, Ph.D. · Department of Information Science, University of Colorado Boulder

What you will build

The page for the 98th Academy Awards (<https://www.oscars.org/oscars/ceremonies/2026>) lists 24 award categories. It has no `<table>`. Each category is a *card*: a name, then one block for each honoree, first the winner and then the nominees. Each honoree block names a person or group and a film (Figure 1). This handout builds the loop from Wednesday's slide one piece at a time and turns the cards into a DataFrame with one row per honoree:

category	honoree_type	entity	film
Actor in a Leading Role	Winner	Michael B. Jordan	Sinners
Actor in a Leading Role	Nominees	Timothée Chalamet	Marty Supreme
... 125 rows in all, from 24 cards			

Each step adds one thing to the step before it. Run each step in your notebook, and compare your output with the output here before you go to the next step.

- Code that you type.
- What Python shows back.
- 1 ... 7 The same seven parts of a card, in every figure and every step.




Figure 1. One card as the browser shows it, September 24, 2026. The numbers match Figure 2 in Step 3.

Step 1 — Ask for the page, and say who you are

```
import requests
from bs4 import BeautifulSoup
import pandas as pd

url = "https://www.oscars.org/oscars/ceremonies/2026"
headers = {"User-Agent": "Web Data Science/v1 brian.keegan@colorado.edu"}
response = requests.get(url, headers=headers)
response.status_code
200
```

The User-Agent header tells the server who is asking: this course and a contact address. Without it, `oscars.org` answers 403 (Forbidden). 200 means that the server sent the page. In your own projects, put your own name and address in the header.

Step 2 — Parse the HTML

```
soup = BeautifulSoup(response.text, "html.parser")
soup.title.get_text()
'The 98th Academy Awards | 2026'
```


`soup` holds the page as a tree of tags that you can search. The title tells you that you have the ceremony page, not an error page.

Step 3 — Find one card in DevTools

Open the page in Chrome, right-click the name MICHAEL B. JORDAN, and select **Inspect**. The Elements panel opens with the `<div>` for that name selected (Figure 2). Each indented line sits inside the nearest less-indented line above it: read up to find the containers around the name, and down to find the parts next to it. Click a triangle to open a line; Figure 2 has the film's `<div class="field">` open.

The class attribute on each line holds several class names, separated by spaces; the line at 5 has three: `field`, `field--name-field-honoree-type`, and `winner`. `oscars.org` is built with Drupal, a content management system that makes most of these names from the site's own field names: `field--name-` plus a field name, and `paragraph--type--` plus a type of block.

These names tell you what a div holds; names such as `field` and `field__item` tell you only that the div is a box.



What each part holds

- 1 the area that holds all 24 cards
- 2 one card (a category)
- 3 the card's name
- 4 one honoree, a winner or a nominee
- 5 Winner or Nominees
- 6 the person or group (the *entity*)
- 7 the film

Figure 2. Chrome DevTools, Elements panel, after **Inspect** on MICHAEL B. JORDAN. September 24, 2026.

Step 4 — Match one class name, then count the matches

```
name_tag = soup.find("div", class_="field--name-field-award-category-oscars")
name_tag["class"]
['field',
 'field--name-field-award-category-oscars',
 'field--type-entity-reference',
 'field--label-hidden',
 'field__item']
```

BeautifulSoup keeps the class attribute as a list of names. `class_="..."` matches a tag when your name is any one item in that list, so you give only the one name that you need. (The underscore is necessary because class is a reserved word in Python.) Now count what each name matches on the whole page. The last item is the whole attribute from the line at 5, copied from DevTools:

```
for name in ["field", "field__item", "paragraph",
             "field--name-field-award-categories", "paragraph--type--award-category",
             "field--name-field-award-category-oscars", "paragraph--type--award-honoree",
             "field--name-field-honoree-type", "field--name-field-award-entities",
             "field--name-field-award-film",
             "field field--name-field-honoree-type winner"]:
    print(len(soup.find_all("div", class_=name)), name)
```

```
693 field
445 field__item
154 paragraph
1 field--name-field-award-categories
24 paragraph--type--award-category
24 field--name-field-award-category-oscars
125 paragraph--type--award-honoree
125 field--name-field-honoree-type
125 field--name-field-award-entities
125 field--name-field-award-film
25 field field--name-field-honoree-type winner
```

The general names match hundreds of divs across the page; the specific names match only what you want. Check each count against the page: it shows 24 category headings, so 24 cards 2 is correct, and 125 honorees 4 with 125 of each field 5 6 7 means one of each per honoree. The whole attribute finds only the 25 winners, because a string with spaces must match the attribute exactly, names and order. Select the one name that describes the content, then check its count.

Step 5 — Go to the area 1, then collect the cards 2

```
area = soup.find("div", class_="field--name-field-award-categories")
categories = area.find_all("div", class_="paragraph--type--award-category")
len(categories)
24
```

`find` returns the first match: one tag. `find_all` returns a list of all the matches. You call `find_all` on `area`, not on `soup`, so it searches only inside the awards area. `categories` is a list of the 24 cards, in the order of the page.

Step 6 — Take one card and read its name 3

```
cat = categories[0]
name_tag = cat.find(class_="field--name-field-award-category-oscar")
name_tag.get_text(strip=True)
'Actor in a Leading Role'
```

Build the steps for one card before you build them for 24. `cat.find(...)` searches only inside this card. This `find` has no "div", so it matches any tag with the class; every field on this page is a div, so both forms work. `get_text(strip=True)` gives the text inside the tag, without the spaces and line breaks around it.

Step 7 — Collect the honorees 4 in that one card

```
honorees = cat.find_all("div", class_="paragraph--type--award-honoree")
len(honorees)
5
```

In Step 4, this class name matched 125 divs on the whole page. Searched from `cat`, it matches the 5 in this card. Where you start a search decides what you get back. `find_all` also looks at every level below `cat`, so you do not step through the two wrapper divs between the card and its honorees (`field--name-field-award-honorees` and `field__item` in Figure 2).

Step 8 — Extract the three fields 5 6 7 from one honoree

```
h = honorees[0]
type_tag = h.find(class_="field--name-field-honoree-type")
entity_tag = h.find(class_="field--name-field-award-entities")
film_tag = h.find(class_="field--name-field-award-film")
print(type_tag.get_text(strip=True))
print(entity_tag.get_text(strip=True))
print(film_tag.get_text(strip=True))
```

```
Winner
Michael B. Jordan
Sinners
```

Each `find` starts from `h`, so it finds only the fields of this honoree. To see what `strip=True` removes, look at the raw text:

```
entity_tag.text
'\nMichael B. Jordan\n'
```

The line breaks come from the HTML around the inner `field__item` div. The type is also a class name (winner or nominee), but the text is what you want in the table.

Step 9 — Make one row for each honoree in the card

```
category_name = name_tag.get_text(strip=True)
rows = []
for h in honorees:
    rows.append({
        "category": category_name,
        "honoree_type": h.find(class_="field--name-field-honoree-type").get_text(strip=True),
        "entity": h.find(class_="field--name-field-award-entities").get_text(strip=True),
        "film": h.find(class_="field--name-field-award-film").get_text(strip=True),
    })
rows[1]
{'category': 'Actor in a Leading Role',
 'honoree_type': 'Nominees',
 'entity': 'Timothée Chalamet',
 'film': 'Marty Supreme'}
```

The three `find` calls from Step 8 go inside the loop, and their text goes into a dictionary: one dictionary per row, keys as column names (`rows[1]` is the second of 5). `category_name`, from Step 6, puts the card's name in each row. The page

shows NOMINEES once, but the HTML has the label in every nominee block: a script hides the repeats in the browser (style="display: none;" in DevTools), and requests does not run scripts, so every row gets its label.

Step 10 — Put the card loop inside a loop over all 24 cards

```
rows = []
for cat in categories:                                # Step 5
    name_tag = cat.find(class_="field--name-field-award-category-oscars") # Step 6
    if name_tag is None:                               # guard
        continue
    category_name = name_tag.get_text(strip=True)      # Step 6
    for h in cat.find_all("div", class_="paragraph--type--award-honoree"): # Step 7
        type_tag = h.find(class_="field--name-field-honoree-type")         # Step 8
        entity_tag = h.find(class_="field--name-field-award-entities")      # Step 8
        film_tag = h.find(class_="field--name-field-award-film")            # Step 8
        rows.append({                                                         # Step 9
            "category": category_name,
            "honoree_type": type_tag.get_text(strip=True) if type_tag else None,
            "entity": entity_tag.get_text(strip=True) if entity_tag else None,
            "film": film_tag.get_text(strip=True) if film_tag else None,
        })
len(rows)
125
```

Two things are new. The outer loop for `cat in categories:` runs Steps 6 to 9 once per card, and `rows = []` moves above it so all 24 cards add to one list. And guards: `find` returns `None` when nothing matches, and `.get_text()` on `None` raises an `AttributeError` that stops the loop. `if name_tag is None: continue` skips a card with no name; ... `if type_tag else None` records a missing field as `None`. No guard fires on the 2026 page; the guards keep the loop running if one card is shaped differently.

Step 11 — Make the DataFrame, then check it against the page

```
df = pd.DataFrame(rows)
df.shape
(125, 4)
df.head()
```

	category	honoree_type	entity	film
0	Actor in a Leading Role	Winner	Michael B. Jordan	Sinners
1	Actor in a Leading Role	Nominees	Timothée Chalamet	Marty Supreme
2	Actor in a Leading Role	Nominees	Leonardo DiCaprio	One Battle after Another
3	Actor in a Leading Role	Nominees	Ethan Hawke	Blue Moon
4	Actor in a Leading Role	Nominees	Wagner Moura	The Secret Agent

```
df["honoree_type"].value_counts()
honoree_type
Nominees    100
Winner       25
Name: count, dtype: int64
```

25 winners for 24 categories looks like a bug in the loop. Look at the winners of one category:

```
df[(df["category"] == "Live Action Short Film") & (df["honoree_type"] == "Winner")]
category      honoree_type  entity      film
70  Live Action Short Film  Winner  Sam A. Davis and Jack Piatt  The Singers
71  Live Action Short Film  Winner  Alexandre Singh and Natalie Musteata  Two People Exchanging Saliva
```

The category ended in a tie (the first Oscars tie since 2013), so the table is correct: check a surprising number against the page before you change your code. Check what each column holds, too. You named the columns, but the site decides what goes in each field: for International Feature Film, entity is the film and film is the country (*Sentimental Value*, Norway); for Music (Original Song), film is the song title.

The whole path

soup → area ① → each card ② → its name ③ and honorees ④ → each honoree's type ⑤, entity ⑥, and film ⑦ → one dictionary per row → `pd.DataFrame(rows)`. If your counts differ, the page has changed: inspect a card again (Figure 2).